

final_project_SK

March 15, 2021

1 Aim of project

1.1 - To make a model to predict the stage of Alzheimers using Convolutional Neural Networks

Data is available at <https://www.kaggle.com/tourist55/alzheimers-dataset-4-class-of-images>

About the data directory. The MRI Images have already been converted to .jpg in the dataset Dataset consists of two files - Training and Testing both containing a total of around ~5000 images each segregated into the severity of Alzheimer's

Classes:

1. MildDemented
2. VeryMildDemented
3. NonDemented
4. ModerateDemeneted

```
[2]: import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
[3]: import tensorflow as tf
from tensorflow.keras.preprocessing import image_dataset_from_directory
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import InputLayer, BatchNormalization, Dropout, Flatten, Dense, Activation, MaxPooling2D, Conv2D, ReLU
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
from keras import optimizers
```

```
[4]: # checking tensorflow version
print(tf.__version__)
```

2.3.0

1.2 Checking data type

The data is stored as

- Alzheimers_s Dataset/ test/ MildDemented/ VeryMildDemented/ NonDemented/ ModerateDemented/ train/ MildDemented/ VeryMildDemented/ NonDemented/ ModerateDemented/

```
[5]: main_data_file_path = '/Users/Susheel/Desktop/BI0F399_final_project/Alzheimer_s_
↳Dataset/'
train_data_file_path = main_data_file_path + 'train/'
test_data_file_path = main_data_file_path + 'test/'
```

```
[7]: ## Parameters
```

```
[6]: BATCH_SIZE = 32
IMAGE_SIZE = [176, 208]
```

```
[7]: ### getting the dataset ready
```

```
[8]: train_dataset = tf.keras.preprocessing.image_dataset_from_directory(
    directory = train_data_file_path,
    image_size = IMAGE_SIZE,
    seed = 1000,
    subset = 'training',
    batch_size = BATCH_SIZE,
    validation_split = 0.2)

validataion_dataset = tf.keras.preprocessing.image_dataset_from_directory(
    directory = train_data_file_path,
    image_size = IMAGE_SIZE,
    seed = 1000,
    subset = 'training',
    batch_size = BATCH_SIZE,
    validation_split = 0.2)

test_dataset = tf.keras.preprocessing.image_dataset_from_directory(
    directory = test_data_file_path,
    image_size = IMAGE_SIZE,
    batch_size = BATCH_SIZE)
```

```
Found 5122 files belonging to 4 classes.
Using 4098 files for training.
Found 5122 files belonging to 4 classes.
Using 4098 files for training.
Found 1279 files belonging to 4 classes.
```

Defining the data classes


```
[9]: class_names = ['MildDementia', 'ModerateDementia', 'NonDementia',
↳'VeryMildDementia']
train_dataset.class_names = class_names
validataion_dataset.class_names = class_names
```

```
NUM_CLASSES = len(class_names)
```

1.3 Data visualization

```
[13]: plt.figure(figsize=(10, 10))
      for images, labels in train_dataset.take(1):
          for i in range(9):
              ax = plt.subplot(3, 3, i + 1)
              plt.imshow(images[i].numpy().astype("uint8"))
              plt.title(train_dataset.class_names[labels[i]])
              plt.axis("off")

      plt.savefig('Example_data_type.png')
```



2 Feature Engineering

Because we are working with categorical and noncontinuous data, we want to convert our model into one-hot encodings. One-hot encodings are a way for the model to understand that we're looking at categorical instead of continuous data. Transforming features so that they'll be more understandable is called feature engineering.

```
[10]: def one_hot_label(image, label):
        label = tf.one_hot(label, NUM_CLASSES)
        return image, label

train_dataset = train_dataset.map(one_hot_label)
validation_dataset = validation_dataset.map(one_hot_label)
test_dataset = test_dataset.map(one_hot_label)
```

3 Model Building

```
[11]: def build_model():

        '''Sequential Model creation'''
        model = Sequential()

        model.add(Conv2D(16, (3,3), padding='same', input_shape =
→ (IMAGE_SIZE[0], IMAGE_SIZE[1], 3), activation='relu'))
        model.add(BatchNormalization())
        model.add(MaxPooling2D(pool_size=(2,2), strides=2, padding = 'same'))

        model.add(Conv2D(32, (3,3), padding='same', activation='relu'))
        model.add(BatchNormalization())
        model.add(MaxPooling2D(pool_size=(2,2), strides=2, padding = 'same'))

        model.add(Conv2D(64, (3,3), padding='same', activation='relu'))
        model.add(BatchNormalization())
        model.add(MaxPooling2D(pool_size=(2,2), strides=2, padding = 'same'))

        model.add(Conv2D(128, (3,3), padding='same', activation='relu'))
        model.add(MaxPooling2D(pool_size=(2,2), strides=2, padding = 'same'))

        model.add(Flatten())
        model.add(Dense(32))
        model.add(ReLU())
        model.add(Dropout(0.25)) # Avoid over-fitting
        model.add(Dense(4))
        model.add(Activation('softmax'))

        return model
```

```
model = build_model()
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 176, 208, 16)	448
batch_normalization (Batch Normalization)	(None, 176, 208, 16)	64
max_pooling2d (MaxPooling2D)	(None, 88, 104, 16)	0
conv2d_1 (Conv2D)	(None, 88, 104, 32)	4640
batch_normalization_1 (Batch Normalization)	(None, 88, 104, 32)	128
max_pooling2d_1 (MaxPooling2D)	(None, 44, 52, 32)	0
conv2d_2 (Conv2D)	(None, 44, 52, 64)	18496
batch_normalization_2 (Batch Normalization)	(None, 44, 52, 64)	256
max_pooling2d_2 (MaxPooling2D)	(None, 22, 26, 64)	0
conv2d_3 (Conv2D)	(None, 22, 26, 128)	73856
max_pooling2d_3 (MaxPooling2D)	(None, 11, 13, 128)	0
flatten (Flatten)	(None, 18304)	0
dense (Dense)	(None, 32)	585760
re_lu (ReLU)	(None, 32)	0
dropout (Dropout)	(None, 32)	0
dense_1 (Dense)	(None, 4)	132
activation (Activation)	(None, 4)	0
Total params: 683,780		
Trainable params: 683,556		
Non-trainable params: 224		

```
[16]: # saving the model summary to a text file
from contextlib import redirect_stdout

with open('modelsummary.txt', 'w') as f:
    with redirect_stdout(f):
        model.summary()
```

```
[17]: METRICS = [tf.keras.metrics.BinaryAccuracy(name='accuracy'),
                 tf.keras.metrics.Precision(name='precision'),
                 tf.keras.metrics.Recall(name='recall'),
                 tf.keras.metrics.AUC(name='auc')]
```

3.1 Hyperparameters

```
[18]: num_epochs = 10
initial_learning_rate = 0.001
# optimizer is ADAM
opt = tf.keras.optimizers.Adam(learning_rate=initial_learning_rate, name='Adam')
```

```
[19]: model.compile(optimizer=opt,
                    loss='categorical_crossentropy',
                    metrics=METRICS)
```

3.2 Defining callbacks

```
[20]: checkpoint_cb = tf.keras.callbacks.ModelCheckpoint("alzheimer_model.h5",
                                                         save_best_only=True)

early_stopping_cb = tf.keras.callbacks.EarlyStopping(patience=10,
                                                         restore_best_weights=True)
```

4 Model fitting

```
[21]: history = model.fit(train_dataset,
                           validation_data=validation_dataset,
                           callbacks=[checkpoint_cb, early_stopping_cb],
                           epochs= num_epochs
                           )
```

Epoch 1/10

129/129 [=====] - 151s 1s/step - loss: 1.4330 - accuracy: 0.7534 - precision: 0.5135 - recall: 0.2557 - auc: 0.7544 - val_loss: 0.9853 - val_accuracy: 0.7953 - val_precision: 0.7127 - val_recall: 0.3038 - val_auc: 0.8142

Epoch 2/10

129/129 [=====] - 148s 1s/step - loss: 0.9891 - accuracy: 0.7862 - precision: 0.6398 - recall: 0.3316 - auc: 0.8088 - val_loss:

0.9576 - val_accuracy: 0.7762 - val_precision: 0.5645 - val_recall: 0.4583 -
val_auc: 0.8258

Epoch 3/10

129/129 [=====] - 158s 1s/step - loss: 0.9426 -
accuracy: 0.7915 - precision: 0.6613 - recall: 0.3402 - auc: 0.8223 - val_loss:
0.8696 - val_accuracy: 0.7981 - val_precision: 0.6469 - val_recall: 0.4239 -
val_auc: 0.8431

Epoch 4/10

129/129 [=====] - 184s 1s/step - loss: 0.8926 -
accuracy: 0.7976 - precision: 0.6800 - recall: 0.3599 - auc: 0.8407 - val_loss:
0.8242 - val_accuracy: 0.8087 - val_precision: 0.7137 - val_recall: 0.3924 -
val_auc: 0.8721

Epoch 5/10

129/129 [=====] - 182s 1s/step - loss: 0.8442 -
accuracy: 0.8049 - precision: 0.7093 - recall: 0.3721 - auc: 0.8585 - val_loss:
0.8034 - val_accuracy: 0.8217 - val_precision: 0.8462 - val_recall: 0.3504 -
val_auc: 0.8982

Epoch 6/10

129/129 [=====] - 183s 1s/step - loss: 0.8051 -
accuracy: 0.8093 - precision: 0.7156 - recall: 0.3936 - auc: 0.8735 - val_loss:
0.8101 - val_accuracy: 0.8236 - val_precision: 0.9235 - val_recall: 0.3209 -
val_auc: 0.8994

Epoch 7/10

129/129 [=====] - 194s 2s/step - loss: 0.7435 -
accuracy: 0.8231 - precision: 0.7613 - recall: 0.4258 - auc: 0.8941 - val_loss:
0.6928 - val_accuracy: 0.8062 - val_precision: 0.6469 - val_recall: 0.4954 -
val_auc: 0.8986

Epoch 8/10

129/129 [=====] - 189s 1s/step - loss: 0.7348 -
accuracy: 0.8297 - precision: 0.7759 - recall: 0.4485 - auc: 0.8985 - val_loss:
0.5820 - val_accuracy: 0.9104 - val_precision: 0.9061 - val_recall: 0.7157 -
val_auc: 0.9503

Epoch 9/10

129/129 [=====] - 189s 1s/step - loss: 0.6601 -
accuracy: 0.8621 - precision: 0.8296 - recall: 0.5642 - auc: 0.9198 - val_loss:
0.5397 - val_accuracy: 0.9289 - val_precision: 0.9511 - val_recall: 0.7545 -
val_auc: 0.9666

Epoch 10/10


129/129 [=====] - 204s 2s/step - loss: 0.5940 -
accuracy: 0.8885 - precision: 0.8914 - recall: 0.6308 - auc: 0.9374 - val_loss:
0.4327 - val_accuracy: 0.9197 - val_precision: 0.9185 - val_recall: 0.7450 -
val_auc: 0.9700

4.1 Model Visualization

```
[22]: fig, ax = plt.subplots(1, 2, figsize=(20, 3))
      ax = ax.ravel()

      for i, met in enumerate(['auc', 'loss']):
          ax[i].plot(history.history[met])
          ax[i].plot(history.history['val_' + met])
          ax[i].set_title('Model {}'.format(met))
          ax[i].set_xlabel('epochs')
          ax[i].set_ylabel(met)
          ax[i].legend(['train', 'val'])

      plt.savefig("Model_Results.png")
```



4.2 Evaluate the Model

```
[23]: scores = model.evaluate(test_dataset)
```

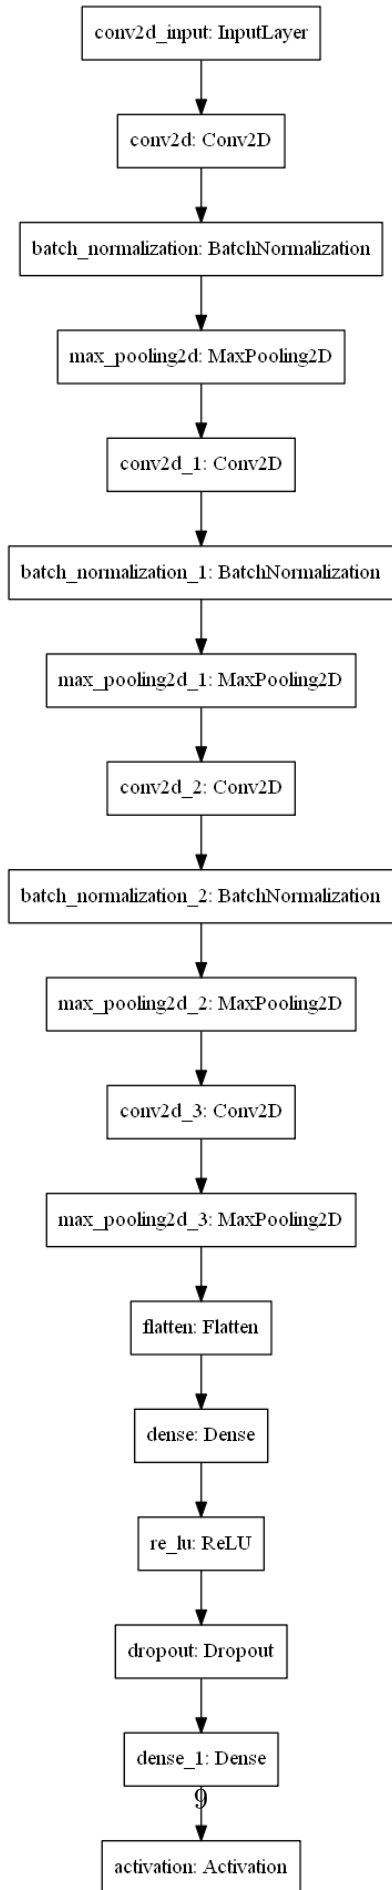
40/40 [=====] - 8s 192ms/step - loss: 1.4601 - accuracy: 0.7961 - precision: 0.5974 - recall: 0.5661 - auc: 0.8329

```
[26]: %%capture cap
      print("Accuracy = ", scores[1])
      print("Precision = ", scores[2])
      print("Recall = ", scores[3])
      print("AUC = ", scores[4])
```

```
[28]: # saving the scores to a file
      with open('modelscores_on_test_dataset.txt', 'w') as f:
          f.write(cap.stdout)
      del cap
```

```
[13]: # Plot the model as a figure
      tf.keras.utils.plot_model(model, to_file='model.png')
```

```
[13]:
```

[]: